

CIFAR-100 Homework_v002

February 23, 2023

You can check and explore further methods applied the tasks we have previously covered from the following link:
https://rodrigob.github.io/are_we_there_yet/build/classification_datasets_results.html#43494641522d313030

Skeleton code is provided from the following tutorial: https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html

Dataset: https://pytorch.org/docs/stable/_modules/torchvision/datasets/cifar.html#CIFAR100

```
[ ]: import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import DataLoader, random_split
import torchvision
import torchvision.transforms as transforms
import numpy as np
import pandas as pd
import cv2
from scipy.ndimage.interpolation import map_coordinates
from scipy.ndimage.filters import gaussian_filter
import matplotlib.pyplot as plt
import math
from timeit import default_timer as timer

print("Versions.....")
print("torch:      ", torch.__version__)
print("torchvision: ", torchvision.__version__)
```

D:\Anaconda3\envs\myEnv001\lib\site-packages\numpy_distributor_init.py:30:

UserWarning: loaded more than 1 DLL from .libs:

D:\Anaconda3\envs\myEnv001\lib\site-packages\numpy\.libs\libopenblas.EL2C6PLE4ZYW3ECEVIV30XXGRN2NRFM2.gfortran-win_amd64.dll

D:\Anaconda3\envs\myEnv001\lib\site-packages\numpy\.libs\libopenblas.PYQHXLVVQ7VESDPUVUADXEVJOBGHJPAY.gfortran-win_amd64.dll

warnings.warn("loaded more than 1 DLL from .libs:")

Versions...

torch: 1.10.1+cu102

torchvision: 0.11.2+cu102

0.0.1 Defining and testing custom image transformation

```
[ ]: # [DONE] TODO: write at least one data transform or augmentation method
      ↪yourself ==> DONE
class ElasticDeformation(object):
    """Elastic deformation used as image augmentation for classification task,
    ↪as per the following paper.
        Simard, Steinkraus and Platt, "Best Practices for Convolutional Neural
    ↪Networks applied to Visual Document Analysis",
        in Proc. of the International Conference on Document Analysis and
    ↪Recognition, 2003.

    Note that calling this class's function will treat its image input
    ↪according to its class.
        If the image is a numpy array, it will be treated like an OpenCV image
    ↪(channels last; H,W,C)
        If the image is a torch.tensor, it will be treated accordingly (channels
    ↪first; C,H,W)

    Args:

        :param alpha: Size of distortion (pixels). typical value is between 1.5 and
    ↪3 x image height
        :param sigma: Variance of the Gaussian distribution with whom the uniform
    ↪random displacement fields will be convolved. Typical value is between 0.01
    ↪and 0.1 image height.
        :param alpha_affine: Size of the affine transformations. Typical value is
    ↪between 0.01 and 0.1 image height.
        :param random_state: State of the RNG for reproducibility of random
    ↪displacement fields. Default is None.
        :param probability: Chance of the transform happening (it is a random
    ↪deformation). Default is 1.0, meaning the transformation will definitely
    ↪happen.
        :param output: Type of output for its callable function. Options are
    ↪'array' for OpenCV output (channels last), or 'tensor' for PyTorch Tensor
    ↪(channels first)
    """
    def __init__(self, alpha, sigma, alpha_affine, random_state=None,
    ↪probability=1.0, output='array'):
        self.alpha = alpha
        self.sigma = sigma
        self.alpha_affine = alpha_affine
        self.random_state = random_state
        self.probability = probability
```

```

self.output = output

def __call__(self, image):
    """Perform elastic transformation on an image using the class
    ↪ parameters for transformation.

    Args:
        :param image (numpy.ndarray or torch.tensor): Input image or tensor
        ↪ of image.

        Note that calling this class's function will treat its image
        ↪ input according to its class.

        If the image is a numpy array, it will be treated like an
        ↪ OpenCV image (channels last; H,W,C)

        If the image is a torch.tensor, it will be treated accordingly
        ↪ (channels first; C,H,W)

    Returns:
        :return image: If output=='array' in constructor, output will be an
        ↪ OpenCV array (channels last), if output=='tensor', a PyTorch tensor
        ↪ (channels first) will be returned.
    """
    if self.random_state is None:
        random_state = np.random.RandomState(None)

    if np.random.uniform() > self.probability:
        if 'tensor' in type(image).__name__.lower():
            if self.output=='tensor':
                return image
            else:
                return image.numpy().transpose((1,2,0))
        else:
            if self.output=='tensor':
                return torch.tensor(image.transpose((2,0,1)))
            else:
                return image

    image = image.numpy().transpose((1,2,0)) if 'tensor' in type(image).
    ↪ __name__.lower() else image

    shape = image.shape
    shape_size = shape[:2]

    # Random affine
    center_square = np.float32(shape_size) // 2
    square_size = min(shape_size) // 3

```

```

        pts1 = np.float32([center_square + square_size,
↪[center_square[0]+square_size, center_square[1]-square_size], center_square,
↪- square_size])
        pts2 = pts1 + random_state.uniform(-self.alpha_affine, self.
↪alpha_affine, size=pts1.shape).astype(np.float32)
        M = cv2.getAffineTransform(pts1, pts2)
        image = cv2.warpAffine(image, M, shape_size[::-1], borderMode=cv2.
↪BORDER_REFLECT_101)

        dx = gaussian_filter((random_state.rand(*shape) * 2 - 1), self.sigma) *
↪self.alpha
        dy = gaussian_filter((random_state.rand(*shape) * 2 - 1), self.sigma) *
↪self.alpha
        dz = np.zeros_like(dx)

        x, y, z = np.meshgrid(np.arange(shape[1]), np.arange(shape[0]), np.
↪arange(shape[2]))
        indices = np.reshape(y+dy, (-1, 1)), np.reshape(x+dx, (-1, 1)), np.
↪reshape(z, (-1, 1))

        out = map_coordinates(image, indices, order=1, mode='reflect').
↪reshape(shape)
        return out if self.output=='array' else torch.tensor(out.
↪transpose((2,0,1)))

    def __repr__(self):
        return self.__class__.__name__+"(\n"+\
            "    alpha = {},\n".format(self.alpha)+\
            "    sigma = {},\n".format(self.sigma)+\
            "    alpha_affine = {},\n".format(self.alpha_affine)+\
            "    random_state = {},\n".format(self.random_state)+\
            "    probability = {},\n".format(self.probability)+\
            "    output = {})".format(self.output)

def elastic_transform_batch(batch, alpha, sigma, alpha_affine,
↪random_state=None, probability=1.0):
    # Perform Elastic Deformation on a whole batch. Here, batch can also be a
↪torch tensor from a dataloader or something.
    distorter = ElasticDeformation(alpha, sigma, alpha_affine, random_state,
↪probability, output='array')
    return np.concatenate([distorter(img) for img in batch]).reshape(batch.
↪shape)

```

```

# Testing the elastic distorter
sampleset = torchvision.datasets.CIFAR100(root='./data/sample', train=False,
↳download=True)
sampledata = sampleset.data[np.random.choice(len(sampleset),10),...]
h,w = sampledata.shape[1:3]
sampledata_deformed = elastic_transform_batch(sampledata, w*1.1, w*0.1, w*0.02)
plt.figure(figsize = (5,25))
i = 0
for idx in range(10):
    i += 1
    plt.subplot(10,2,i)
    plt.imshow(sampledata[idx,...])
    plt.axis('off')
    if i == 1: plt.title('Original')
    i += 1
    plt.subplot(10,2,i)
    plt.imshow(sampledata_deformed[idx,...])
    if i == 2: plt.title('Deformed')
    plt.axis('off')
plt.show()










```

Files already downloaded and verified

Original



Deformed



0.0.2 Preliminaries and Data Loaders

```
[ ]: # [DONE] TODO: You may consider applying more transform such as data_
      ↪ augmentation methods, etc.
      # TODO: You may consider hyperparameter optimization: in this cell, we have_
      ↪ batch_size!
      # [DONE] TODO: use the previously defined data transform/augmentation method in_
      ↪ the following transform.

      # Global constants
      ELASTIC_ALPHA = 32*1.1
      ELASTIC_SIGMA = 32*0.1
      ELASTIC_ALPHA_AFFINE = 32*0.02
      ELASTIC_PROBABILITY = 0.5
      P_VERT_FLIP = 0.5
      P_HORZ_FLIP = 0.5
      JITTER_BRIGHTNESS = 0.1
      JITTER_SATURATION = 0.1
      JITTER_CONTRAST = 0.2
      AFFINE_ROTATION = 50
      AFFINE_TRANSLATION = (0.1, 0.1)
      AFFINE_SCALE = (0.5, 1.5)
      AFFINE_SHEAR = 20
      NORM_MEAN = (0.5, 0.5, 0.5)
      NORM_STD = (0.5, 0.5, 0.5)
      VALIDATION_DATA = 0.2
      NUM_HPO_TRIALS = 10
      HPO_TIMEOUT_SEC = 3600 * 5
      NUM_HPO_JOBS = 1

      # Hyperparameters
      hparams = {
          'input_image_size': [32, 32],
          'batch_size': 256,
          'epochs': 10,
          'lr': 0.001,
          'loss_function': nn.CrossEntropyLoss(),
          'optimizer': 'adam',
          'optimizer_params': None,
          'num_conv_blocks': 3,
          'num_dense_layers': 3,
          'conv_kernel_size': 5,
          'pool_kernel_size': 2,
          'conv_padding': 0,
```

```

    'pool_padding':0,
    'conv_activations':'relu',
    'dense_activations':'relu',
    'conv_batchnorm':'before',
    'dense_batchnorm':'before',
    'conv_stride':1,
    'pool_stride':1,
    'dense_sizes':[256,128],
    'conv_dropout':0.2,
    'dense_dropout':0.2,
    'L2':0.001,
    'num_classes':100}

transform = transforms.Compose([
    transforms.ToTensor(),
    ElasticDeformation(ELASTIC_ALPHA, ELASTIC_SIGMA, ELASTIC_ALPHA_AFFINE,
    ↪random_state=None, probability=ELASTIC_PROBABILITY, output='tensor'),
    transforms.RandomVerticalFlip(P_VERT_FLIP),
    transforms.RandomHorizontalFlip(P_HORZ_FLIP),
    transforms.ColorJitter(brightness=JITTER_BRIGHTNESS,
    ↪contrast=JITTER_CONTRAST, saturation=JITTER_SATURATION),
    transforms.RandomAffine(AFFINE_ROTATION, AFFINE_TRANSLATION,
    ↪AFFINE_SCALE, AFFINE_SHEAR),
    transforms.Normalize(NORM_MEAN, NORM_STD)])

allset = torchvision.datasets.CIFAR100(root='./data/CIFAR100', train=True,
    ↪download=True, transform=transform)
num_all_data = len(allset.data)
num_val_data, num_train_data = int(VALIDATION_DATA*num_all_data),
    ↪int((1-VALIDATION_DATA)*num_all_data)
(trainset, valset) = random_split(allset, (num_train_data, num_val_data),
    ↪generator=torch.Generator().manual_seed(42))
trainloader = DataLoader(trainset, batch_size=hparams['batch_size'],
    ↪shuffle=True, num_workers=0)
validloader = DataLoader(valset, batch_size=hparams['batch_size'],
    ↪shuffle=True, num_workers=0)
testset = torchvision.datasets.CIFAR100(root='./data/CIFAR100', train=False,
    ↪download=True, transform=transform)
testloader = DataLoader(testset, batch_size=hparams['batch_size'],
    ↪shuffle=False, num_workers=0)

classes = ('apples', 'aquarium fish', 'baby', 'bear', 'beaver', 'bed', 'bee',
    ↪'beetle', 'bicycle', 'bottles',

```



```

        'bowls', 'boy', 'bridge', 'bus', 'butterfly', 'camel', 'cans',
↪ 'castle', 'caterpillar', 'cattle', 'chair',
        'chimpanzee', 'clock', 'cloud', 'cockroach', 'computer keyboard',
↪ 'couch', 'crab', 'crocodile',
        'cups', 'dinosaur', 'dolphin', 'elephant', 'flatfish', 'forest',
↪ 'fox', 'girl', 'hamster', 'house',
        'kangaroo', 'lamp', 'lawn-mower', 'leopard', 'lion', 'lizard',
↪ 'lobster', 'man', 'maple', 'motorcycle',
        'mountain', 'mouse', 'mushrooms', 'oak', 'oranges', 'orchids',
↪ 'otter', 'palm', 'pears', 'pickup truck',
        'pine', 'plain', 'plates', 'poppies', 'porcupine', 'possum',
↪ 'rabbit', 'raccoon', 'ray', 'road', 'rocket',
        'roses', 'sea', 'seal', 'shark', 'shrew', 'skunk', 'skyscraper',
↪ 'snail', 'snake', 'spider', 'squirrel',
        'streetcar', 'sunflowers', 'sweet peppers', 'table', 'tank',
↪ 'telephone', 'television', 'tiger', 'tractor',
        'train', 'trout', 'tulips', 'turtle', 'wardrobe', 'whale', 'willow',
↪ 'wolf', 'woman', 'worm')

```

Files already downloaded and verified

Files already downloaded and verified

0.0.3 Visualizing some normalized images

```

[ ]: def imshow(img):
    img = img / 2 + 0.5      # unnormalize
    npimg = img.numpy()
    plt.figure(figsize=(12,3))
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# get some random training images
dataiter = iter(trainloader)


images, labels = dataiter.next()
#images, labels = testloader[0]
print("Shape of images: ", images.shape)
print("Shape of labels: ", labels.shape)
print(labels)
# show images
imshow(torchvision.utils.make_grid(images[:16]))
# print labels
print(', '.join('%5s' % classes[labels[j]] for j in range(16)))

```

Shape of images: torch.Size([1024, 3, 32, 32])

Shape of labels: torch.Size([1024])

tensor([49, 90, 8, ..., 82, 9, 84])



mountain, train, bicycle, plain, squirrel, mountain, poppies, trout, cups,
beaver, crab, shark, rabbit, motorcycle, willow, castle

0.0.4 Constructing CNN model

```
[ ]: # Base model, not to be used here
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 16, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(16, 32, 5)
        self.fc1 = nn.Linear(32 * 5 * 5, 200)
        self.fc2 = nn.Linear(200, 128)
        self.fc3 = nn.Linear(128, 100)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 32 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

```
[ ]: # [DONE] (Optional but highly recommended) TODO: You can define your own neural_
      ↪ network to create a better performing model!
class ConvNet(nn.Module):
    def __init__(self, hparams=dict):
        super(ConvNet, self).__init__()
        self.hparams = hparams
        self.layers = []
        self._img_size = self.hparams['input_image_size']
        self._h, self._w = self._img_size
        self._img_size_list = [[3, self._h, self._w]]
```

```

        self._actdict = {'relu':nn.ReLU(), 'leakyrelu':nn.LeakyReLU(0.1),
        ↪ 'sigmoid':nn.Sigmoid(), 'tanh':nn.Tanh()})

    # Constructing the encoder (feature extractor)
    in_channels = 3
    out_channels = 16
    for i in range(self.hparams['num_conv_blocks']):
        self.layers.append(nn.Conv2d(in_channels, out_channels, self.
        ↪ hparams['conv_kernel_size'], padding=self.hparams['conv_padding'],
        ↪ stride=self.hparams['conv_stride']))
        self._update_image_size(out_channels, 'conv')
        if self.hparams['conv_batchnorm']=='before':
            self.layers.append(nn.BatchNorm2d(out_channels))
        self.layers.append(self._actdict[self.hparams['conv_activations']])
        if self.hparams['conv_batchnorm']=='after':
            self.layers.append(nn.BatchNorm2d(out_channels))
        self.layers.append(nn.MaxPool2d(self.hparams['pool_kernel_size'],
        ↪ padding=self.hparams['pool_padding'], stride=self.hparams['pool_stride']))
        self._update_image_size(out_channels, 'pool')
        if self.hparams.get('conv_dropout'):
            self.layers.append(nn.Dropout2d(self.hparams['conv_dropout']))
        if i < self.hparams['num_conv_blocks'] - 1:
            in_channels = out_channels
            out_channels *= 2

    # Flattening (Image embedding)
    self.layers.append(nn.Flatten())

    # Constructing the decoder (classifier)
    if self.hparams.get('dense_sizes'):
        size_vec = self.hparams.get('dense_sizes')
    else:
        size_vec = [0] * self.hparams['num_dense_layers']
        size_vec[-1] = np.power(2, np.ceil(np.log2(self.
        ↪ hparams['num_classes'])))
        for i in range(self.hparams['num_dense_layers']-2, -1, -1):
            size_vec[i] = size_vec[i+1] * 2
    in_size = out_channels*self._h*self._w
    out_size = size_vec[0]
    for i in range(len(size_vec)):
        self.layers.append(nn.Linear(in_size, out_size))
        if self.hparams['dense_batchnorm'] == 'before':
            self.layers.append(nn.BatchNorm1d(out_size))
        self.layers.append(self._actdict[self.hparams['dense_activations']])
        if self.hparams['dense_batchnorm'] == 'after':
            self.layers.append(nn.BatchNorm1d(out_size))
        if self.hparams.get('dense_dropout'):

```

```

        self.layers.append(nn.Dropout(self.hparams['dense_dropout']))
    if i < len(size_vec) - 1:
        in_size = out_size
        out_size = size_vec[i+1]
    self.layers.append(nn.Linear(out_size, self.hparams['num_classes']))

    # Constructing model
    self.net = nn.Sequential(*self.layers)

def _calc_size(self, size_in:int, padding:int, kernel_size:int, stride:int):
    if padding == 'valid':
        padding=0
    if padding=='same':
        return size_in
    else:
        return math.floor((size_in + 2*padding - (kernel_size-1) - 1)/
↪stride + 1)

def _update_image_size(self, out_channels, ops:str='conv'):
    (self._h, self._w) = (self._calc_size(sz, self.hparams[ops+'_padding'],
↪self.hparams[ops+'_kernel_size'], self.hparams[ops+'_stride']) for sz in
↪(self._h,self._w))
    self._img_size = (self._h, self._w)
    #print("new size: ",self._img_size)
    self._img_size_list.append([out_channels, self._h, self._w])

def forward(self, x:torch.Tensor):
    return self.net(x)

# Instantiation
inzvaNet = ConvNet(hparams)
print("Built network is:")
print(inzvaNet)
print("Image sizes as they go through the convolution blocks:")
print(np.array(inzvaNet._img_size_list))

```

Built network is:

```

ConvNet(
  (net): Sequential(
    (0): Conv2d(3, 16, kernel_size=(5, 5), stride=(1, 1))
    (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (2): ReLU()
    (3): MaxPool2d(kernel_size=2, stride=1, padding=0, dilation=1,

```

```

ceil_mode=False)
    (4): Dropout2d(p=0.2, inplace=False)
    (5): Conv2d(16, 32, kernel_size=(5, 5), stride=(1, 1))
    (6): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (7): ReLU()
    (8): MaxPool2d(kernel_size=2, stride=1, padding=0, dilation=1,
ceil_mode=False)
    (9): Dropout2d(p=0.2, inplace=False)
    (10): Conv2d(32, 64, kernel_size=(5, 5), stride=(1, 1))
    (11): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (12): ReLU()
    (13): MaxPool2d(kernel_size=2, stride=1, padding=0, dilation=1,
ceil_mode=False)
    (14): Dropout2d(p=0.2, inplace=False)
    (15): Flatten(start_dim=1, end_dim=-1)
    (16): Linear(in_features=18496, out_features=256, bias=True)
    (17): BatchNorm1d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (18): ReLU()
    (19): Dropout(p=0.2, inplace=False)
    (20): Linear(in_features=256, out_features=128, bias=True)
    (21): BatchNorm1d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (22): ReLU()
    (23): Dropout(p=0.2, inplace=False)
    (24): Linear(in_features=128, out_features=100, bias=True)
)
)

```

Image sizes as they go through the convolution blocks:

```

[[ 3 32 32]
 [16 28 28]
 [16 27 27]
 [32 23 23]
 [32 22 22]
 [64 18 18]
 [64 17 17]]

```

0.0.5 Checking GPU availability

```

[ ]: # if you want to train on GPU:
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)

```

```
cuda:0
```

0.0.6 Instantiating CNN and defining training procedure

```
[ ]: inzvaNet.to(device)
inzvaNet.train()

# base optimizer with following parameters:
criterion = hparams['loss_function']
if hparams['optimizer']=='adam':
    if hparams.get('optimizer_params'):
        optimizer = optim.Adam(inzvaNet.parameters(), lr=hparams['lr'],
        ↪**hparams['optimizer_params'])
    else:
        optimizer = optim.Adam(inzvaNet.parameters(), lr=hparams['lr'])
elif hparams['optimizer']=='sgd':
    if hparams.get('optimizer_params'):
        optimizer = optim.SGD(inzvaNet.parameters(), lr=hparams['lr'],
        ↪**hparams['optimizer_params'])
    else:
        optimizer = optim.SGD(inzvaNet.parameters(), lr=hparams['lr'])
else:
    raise ValueError("Sorry, only 'adam' and 'sgd' are supported for now.")

# optimizer = optim.SGD(inzvaNet.parameters(), lr=0.001, momentum=0.9)
# TODO: play with hyperparameters and chosen methods to achieve higher accuracy!
↪ You can apply grid or random search.
```

0.0.7 Training

```
[ ]: # Training loop
print('training starts!')
num_training_batches = len(trainloader)
num_validation_batches = len(validloader)
num_testing_batches = len(testloader)

tStart = timer()
for epoch in range(hparams["epochs"]):

    epoch_loss_training = 0.0
    epoch_loss_validation = 0.0

    inzvaNet.train()
    for i, data in enumerate(trainloader):
        images, labels = data[0].to(device), data[1].to(device)
        optimizer.zero_grad()
        predictions = inzvaNet(images)
        loss = criterion(predictions, labels)
        loss.backward()
```

```

optimizer.step()
epoch_loss_training += loss.item()

inzvaNet.eval()
with torch.no_grad():
    for i, data in enumerate(validloader):
        images, labels = data[0].to(device), data[1].to(device)
        predictions = inzvaNet(images)
        loss = criterion(predictions, labels)
        epoch_loss_validation += loss.item()

    print("Epoch: %2d\t|Training Loss: %.3f\t|Validation Loss: %.3f" %
    ↪(epoch + 1,
        epoch_loss_training / num_training_batches,
        epoch_loss_validation / num_validation_batches))

tFinish = timer()
print('Finished Training.')
print("Training process took %.2f seconds"%(tFinish-tStart))
print("Saving model...")
try:
    torch.save(inzvaNet, "ConvNet.pth")
except Exception as e:
    print(e)
    print("Failed to save the model.")
print("Done.")

```

training starts!

Epoch: 1	Training Loss: 4.454	Validation Loss: 4.230
Epoch: 2	Training Loss: 4.253	Validation Loss: 4.105
Epoch: 3	Training Loss: 4.166	Validation Loss: 4.021
Epoch: 4	Training Loss: 4.112	Validation Loss: 3.967
Epoch: 5	Training Loss: 4.058	Validation Loss: 3.917
Epoch: 6	Training Loss: 4.019	Validation Loss: 3.857
Epoch: 7	Training Loss: 3.985	Validation Loss: 3.860
Epoch: 8	Training Loss: 3.957	Validation Loss: 3.836
Epoch: 9	Training Loss: 3.921	Validation Loss: 3.783
Epoch: 10	Training Loss: 3.908	Validation Loss: 3.780

Finished Training.

Training process took 842.76 seconds.

Saving model...

Done.

```

[ ]: dataiter = iter(testloader)
images, labels = dataiter.next()
images, labels = images[:8].to(device), labels[:8].to(device)

```

```

outputs = inzvaNet(images)
_, predicted = torch.max(outputs.data, 1)
# print images
imshow(torchvision.utils.make_grid(images.cpu()))
print('GroundTruth: ', ' ', '.join('%10s' % classes[labels[j]] for j in range(8)))
print('Predicted:   ', ' ', '.join('%10s' % classes[predicted[j]] for j in
↪range(8)))

```



```

GroundTruth:  mountain,  flatfish,          seal,  mushrooms,          sea,
tulips,        camel,  butterfly
Predicted:      cloud,   leopard,          train,          wolf,    apples,
tiger,    leopard,          skunk

```

```

[ ]: # test on all test data
inzvaNet.eval()
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = inzvaNet(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print("correct predictions: ", correct)
print("total number of data: ", total)
print('Accuracy of the network on the 10000 test images: %d %%' % (100 *
↪correct / total))

```

```

correct predictions: 1229
total number of data: 10000
Accuracy of the network on the 10000 test images: 12 %

```

```

[ ]: # class-wise accuracy
class_correct = list(0. for i in range(100))
class_total = list(0. for i in range(100))
with torch.no_grad():

```



```

for data in testloader:
    images, labels = data
    images, labels = images.to(device), labels.to(device)
    outputs = inzvaNet(images)
    _, predicted = torch.max(outputs, 1)
    c = (predicted == labels).squeeze()
    for i in range(len(labels)):
        label = labels[i]
        class_correct[label] += c[i].item()
        class_total[label] += 1

print("Accuracy of classes:")
for i in range(100):
    print('%3d- %20s : %5.3f %%' % (
        i+1, classes[i], 100 * class_correct[i] / class_total[i]))

```

Accuracy of classes:

```

1-          apples : 23.000 %
2-    aquarium fish : 9.000 %
3-          baby : 10.000 %
4-          bear : 9.000 %
5-        beaver : 12.000 %
6-          bed : 1.000 %
7-          bee : 4.000 %
8-        beetle : 13.000 %
9-        bicycle : 2.000 %
10-        bottles : 0.000 %
11-        bowls : 0.000 %
12-         boy : 0.000 %
13-        bridge : 4.000 %
14-         bus : 7.000 %
15-    butterfly : 6.000 %
16-        camel : 8.000 %
17-         cans : 0.000 %
18-        castle : 26.000 %
19-    caterpillar : 28.000 %
20-        cattle : 2.000 %
21-         chair : 13.000 %
22-    chimpanzee : 36.000 %
23-         clock : 0.000 %
24-         cloud : 52.000 %
25-    cockroach : 44.000 %
26- computer keyboard : 0.000 %
27-         couch : 1.000 %
28-         crab : 2.000 %
29-    crocodile : 0.000 %
30-         cups : 4.000 %
31-    dinosaur : 13.000 %

```

32-	dolphin	: 14.000 %
33-	elephant	: 0.000 %
34-	flatfish	: 6.000 %
35-	forest	: 4.000 %
36-	fox	: 6.000 %
37-	girl	: 22.000 %
38-	hamster	: 4.000 %
39-	house	: 6.000 %
40-	kangaroo	: 0.000 %
41-	lamp	: 0.000 %
42-	lawn-mower	: 6.000 %
43-	leopard	: 31.000 %
44-	lion	: 34.000 %
45-	lizard	: 1.000 %
46-	lobster	: 5.000 %
47-	man	: 10.000 %
48-	maple	: 2.000 %
49-	motorcycle	: 7.000 %
50-	mountain	: 24.000 %
51-	mouse	: 0.000 %
52-	mushrooms	: 4.000 %
53-	oak	: 66.000 %
54-	oranges	: 57.000 %
55-	orchids	: 17.000 %
56-	otter	: 0.000 %
57-	palm	: 3.000 %
58-	pears	: 10.000 %
59-	pickup truck	: 1.000 %
60-	pine	: 4.000 %
61-	plain	: 18.000 %
62-	plates	: 20.000 %
63-	poppies	: 39.000 %
64-	porcupine	: 4.000 %
65-	possum	: 4.000 %
66-	rabbit	: 0.000 %
67-	raccoon	: 9.000 %
68-	ray	: 14.000 %
69-	road	: 31.000 %
70-	rocket	: 5.000 %
71-	roses	: 32.000 %
72-	sea	: 13.000 %
73-	seal	: 1.000 %
74-	shark	: 41.000 %
75-	shrew	: 1.000 %
76-	skunk	: 32.000 %
77-	skyscraper	: 4.000 %
78-	snail	: 3.000 %
79-	snake	: 1.000 %

```

80-         spider : 3.000 %
81-         squirrel : 0.000 %
82-         streetcar : 1.000 %
83-         sunflowers : 73.000 %
84-         sweet peppers : 3.000 %
85-         table : 0.000 %
86-         tank : 11.000 %
87-         telephone : 22.000 %
88-         television : 1.000 %
89-         tiger : 4.000 %
90-         tractor : 7.000 %
91-         train : 5.000 %
92-         trout : 1.000 %
93-         tulips : 5.000 %
94-         turtle : 2.000 %
95-         wardrobe : 33.000 %
96-         whale : 36.000 %
97-         willow : 39.000 %
98-         wolf : 11.000 %
99-         woman : 5.000 %
100-        worm : 1.000 %

```

0.0.8 Hyper-Parameter Optimization

```

[ ]: # installing optuna
import subprocess
command = "pip install --no-input optuna"
process = subprocess.run(command, stdout=subprocess.PIPE, stderr=subprocess.
    ↳PIPE, universal_newlines=True, shell=True)
stdout_raw = process.stdout.strip().upper()
stderr_raw = process.stderr.strip().upper()
if len(stderr_raw) > 0:
    print(stderr_raw)
    raise RuntimeError("ERROR: Could not execute shell command '{}'.
    ↳format(command))
print(stdout_raw)
print("Done.")

```

```

REQUIREMENT ALREADY SATISFIED: OPTUNA IN D:\ANACONDA3\ENVS\MYENV001\LIB\SITE-
PACKAGES (2.10.0)
REQUIREMENT ALREADY SATISFIED: SQLALCHEMY>=1.1.0 IN
D:\ANACONDA3\ENVS\MYENV001\LIB\SITE-PACKAGES (FROM OPTUNA) (1.4.29)
REQUIREMENT ALREADY SATISFIED: SCIPY!=1.4.0 IN
D:\ANACONDA3\ENVS\MYENV001\LIB\SITE-PACKAGES (FROM OPTUNA) (1.5.2)
REQUIREMENT ALREADY SATISFIED: COLORLOG IN D:\ANACONDA3\ENVS\MYENV001\LIB\SITE-
PACKAGES (FROM OPTUNA) (6.6.0)
REQUIREMENT ALREADY SATISFIED: TQDM IN D:\ANACONDA3\ENVS\MYENV001\LIB\SITE-
PACKAGES (FROM OPTUNA) (4.62.3)

```

REQUIREMENT ALREADY SATISFIED: CMAES>=0.8.2 IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM OPTUNA) (0.8.2)
 REQUIREMENT ALREADY SATISFIED: ALEMBIC IN D:\ANACONDA3\ENV\MYENV001\LIB\SITE-
 PACKAGES (FROM OPTUNA) (1.7.5)
 REQUIREMENT ALREADY SATISFIED: PYYAML IN D:\ANACONDA3\ENV\MYENV001\LIB\SITE-
 PACKAGES (FROM OPTUNA) (6.0)
 REQUIREMENT ALREADY SATISFIED: PACKAGING>=20.0 IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM OPTUNA) (21.3)
 REQUIREMENT ALREADY SATISFIED: CLIFF IN D:\ANACONDA3\ENV\MYENV001\LIB\SITE-
 PACKAGES (FROM OPTUNA) (3.10.0)
 REQUIREMENT ALREADY SATISFIED: NUMPY IN D:\ANACONDA3\ENV\MYENV001\LIB\SITE-
 PACKAGES (FROM OPTUNA) (1.22.1)
 REQUIREMENT ALREADY SATISFIED: PYPARSING!=3.0.5,>=2.0.2 IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM PACKAGING>=20.0->OPTUNA)
 (3.0.6)
 REQUIREMENT ALREADY SATISFIED: GREENLET!=0.4.17 IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM SQLALCHEMY>=1.1.0->OPTUNA)
 (1.1.2)
 REQUIREMENT ALREADY SATISFIED: IMPORTLIB-METADATA IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM ALEMBIC->OPTUNA) (4.10.0)
 REQUIREMENT ALREADY SATISFIED: IMPORTLIB-RESOURCES IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM ALEMBIC->OPTUNA) (5.4.0)
 REQUIREMENT ALREADY SATISFIED: MAKO IN D:\ANACONDA3\ENV\MYENV001\LIB\SITE-
 PACKAGES (FROM ALEMBIC->OPTUNA) (1.1.6)
 REQUIREMENT ALREADY SATISFIED: CMD2>=1.0.0 IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM CLIFF->OPTUNA) (2.3.3)
 REQUIREMENT ALREADY SATISFIED: PBR!=2.1.0,>=2.0.0 IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM CLIFF->OPTUNA) (5.8.0)
 REQUIREMENT ALREADY SATISFIED: STEVEDORE>=2.0.1 IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM CLIFF->OPTUNA) (3.5.0)
 REQUIREMENT ALREADY SATISFIED: PRETTYTABLE>=0.7.2 IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM CLIFF->OPTUNA) (2.5.0)
 REQUIREMENT ALREADY SATISFIED: AUTOPAGE>=0.4.0 IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM CLIFF->OPTUNA) (0.4.0)
 REQUIREMENT ALREADY SATISFIED: COLORAMA IN D:\ANACONDA3\ENV\MYENV001\LIB\SITE-
 PACKAGES (FROM COLORLOG->OPTUNA) (0.4.4)
 REQUIREMENT ALREADY SATISFIED: PYREADLINE3 IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM CMD2>=1.0.0->CLIFF->OPTUNA)
 (3.3)
 REQUIREMENT ALREADY SATISFIED: ATTRS>=16.3.0 IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM CMD2>=1.0.0->CLIFF->OPTUNA)
 (21.2.0)
 REQUIREMENT ALREADY SATISFIED: WCWIDTH>=0.1.7 IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM CMD2>=1.0.0->CLIFF->OPTUNA)
 (0.2.5)
 REQUIREMENT ALREADY SATISFIED: PYPERCLIP>=1.6 IN
 D:\ANACONDA3\ENV\MYENV001\LIB\SITE-PACKAGES (FROM CMD2>=1.0.0->CLIFF->OPTUNA)
 (1.8.2)

REQUIREMENT ALREADY SATISFIED: ZIPP>=0.5 IN D:\ANACONDA3\ENVS\MYENV001\LIB\SITE-
 PACKAGES (FROM IMPORTLIB-METADATA->ALEMBIC->OPTUNA) (3.7.0)
 REQUIREMENT ALREADY SATISFIED: MARKUPSATE>=0.9.2 IN
 D:\ANACONDA3\ENVS\MYENV001\LIB\SITE-PACKAGES (FROM MAKO->ALEMBIC->OPTUNA)
 (2.0.1)
 Done.

```

[ ]: import optuna
    from optuna.samplers import TPESampler

[ ]: # Function that runs through training and returns the validation loss to be
    ↪ optimized
    def objective(trial):

        # Cleaning up GPU
        torch.cuda.empty_cache()

        # Defining hyperparameter space
        hparams = {
            'input_image_size': [32, 32],
            'batch_size': 256,
            'epochs': 10,
            'lr': 0.001,
            'loss_function': nn.CrossEntropyLoss(),
            'optimizer': trial.suggest_categorical('optimizer', ['adam', 'sgd']),
            'optimizer_params': None,
            'num_conv_blocks': trial.suggest_int("num_conv_blocks", 2, 5),
            'num_dense_layers': trial.suggest_int("num_dense_layers", 1, 5),
            'conv_kernel_size': trial.suggest_int("conv_kernel_size", 2, 10),
            'pool_kernel_size': trial.suggest_int("pool_kernel_size", 2, 5),
            'conv_padding': 'same',
            'pool_padding': 0,
            'conv_activations': trial.suggest_categorical("conv_activations", ['relu',
    ↪ 'leakyrelu', 'sigmoid', 'tanh']),
            'dense_activations': trial.suggest_categorical("dense_activations", ['relu',
    ↪ 'leakyrelu', 'sigmoid', 'tanh']),
            'conv_batchnorm': trial.suggest_categorical("conv_batchnorm", ["before",
    ↪ "after"]),
            'dense_batchnorm': trial.suggest_categorical("dense_batchnorm", ["before",
    ↪ "after"]),
            'conv_stride': 1,
            'pool_stride': 1,
            'conv_dropout': trial.suggest_float("conv_dropout", 0.1, 0.5),
            'dense_dropout': trial.suggest_float("dense_dropout", 0.1, 0.5),
            'L2': 0.001,
            'num_classes': 100}
  
```

```

# Defining hyperparameter conditions
hparams['pool_padding'] = hparams['pool_kernel_size']//2
if hparams['optimizer'] == 'sgd':
    hparams['optimizer_params'] = {'momentum':0.9}

# Generate data
allset = torchvision.datasets.CIFAR100(root='./data/CIFAR100', train=True,
download=True, transform=transform)
num_all_data = len(allset.data)
num_val_data, num_train_data = int(VALIDATION_DATA*num_all_data),
int((1-VALIDATION_DATA)*num_all_data)
(trainset, valset) = random_split(allset, (num_train_data, num_val_data),
generator=torch.Generator().manual_seed(42))
trainloader = DataLoader(trainset, batch_size=hparams['batch_size'],
shuffle=True, num_workers=0)
validloader = DataLoader(valset, batch_size=hparams['batch_size'],
shuffle=True, num_workers=0)

# Build and train the model as we did before
inzvaNet = ConvNet(hparams)
inzvaNet.to(device)
inzvaNet.train()
criterion = hparams['loss_function']
if hparams['optimizer']=='adam':
    if hparams.get('optimizer_params'):
        optimizer = optim.Adam(inzvaNet.parameters(), lr=hparams['lr'],
**hparams['optimizer_params'])
    else:
        optimizer = optim.Adam(inzvaNet.parameters(), lr=hparams['lr'])
elif hparams['optimizer']=='sgd':
    if hparams.get('optimizer_params'):
        optimizer = optim.SGD(inzvaNet.parameters(), lr=hparams['lr'],
**hparams['optimizer_params'])
    else:
        optimizer = optim.SGD(inzvaNet.parameters(), lr=hparams['lr'])
num_training_batches = len(trainloader)
num_validation_batches = len(validloader)

for epoch in range(hparams["epochs"]):

    epoch_loss_training = 0.0
    epoch_loss_validation = 0.0

    inzvaNet.train()
    for i, data in enumerate(trainloader):

```

```

        images, labels = data[0].to(device), data[1].to(device)
        optimizer.zero_grad()
        predictions = inzvaNet(images)
        loss = criterion(predictions, labels)
        loss.backward()
        optimizer.step()
        epoch_loss_training += loss.item()

    inzvaNet.eval()
    with torch.no_grad():
        for i, data in enumerate(validloader):
            images, labels = data[0].to(device), data[1].to(device)
            predictions = inzvaNet(images)
            loss = criterion(predictions, labels)
            epoch_loss_validation += loss.item()

    print("Epoch: %2d\t|Training Loss: %.3f\t|Validation Loss: %.3f" %
    ↪(epoch + 1,
        epoch_loss_training / num_training_batches,
        epoch_loss_validation / num_validation_batches))

    return epoch_loss_validation

```

```

[ ]: # Create HPO study
study = optuna.create_study(sampler=TPESampler(), direction="minimize")

# Optimize HPO study
tStart = timer()
study.optimize(objective, n_trials=NUM_HPO_TRIALS, timeout=HPO_TIMEOUT_SEC,
    ↪n_jobs=NUM_HPO_JOBS, catch=(Exception,))
tFinish = timer()

# Report
print("HPO process took %d seconds."%(tFinish-tStart))
print("Number of completed trials: {}".format(len(study.trials)))
print("Best trial:")
trial = study.best_trial
print(trial)
print("\tBest Validation Loss: {}".format(trial.value))
print("\tBest Params: ")
for key, value in trial.params.items():
    print("    {}: {}".format(key, value))

```

[I 2022-01-24 03:21:21,604] A new study created in memory with name:
no-name-1f2b7940-f244-4b78-be45-831a171969cd

Files already downloaded and verified

D:\Anaconda3\envs\myEnv001\lib\site-packages\torch\nn\modules\conv.py:442:

UserWarning: Using padding='same' with even kernel lengths and odd dilation may require a zero-padded copy of the input be created (Triggered internally at ..\aten\src\ATen\native\Convolution.cpp:647.)

```
return F.conv2d(input, weight, bias, self.stride,
```

Epoch:	1	Training Loss:	4.679	Validation Loss:	4.425
Epoch:	2	Training Loss:	4.527	Validation Loss:	4.337
Epoch:	3	Training Loss:	4.447	Validation Loss:	4.281
Epoch:	4	Training Loss:	4.389	Validation Loss:	4.231
Epoch:	5	Training Loss:	4.342	Validation Loss:	4.189
Epoch:	6	Training Loss:	4.290	Validation Loss:	4.148
Epoch:	7	Training Loss:	4.260	Validation Loss:	4.120
Epoch:	8	Training Loss:	4.226	Validation Loss:	4.071
Epoch:	9	Training Loss:	4.197	Validation Loss:	4.063

[I 2022-01-24 03:38:16,548] Trial 0 finished with value: 160.976891040802 and parameters: {'optimizer': 'sgd', 'num_conv_blocks': 3, 'num_dense_layers': 4, 'conv_kernel_size': 4, 'pool_kernel_size': 5, 'conv_activations': 'relu', 'dense_activations': 'leakyrelu', 'conv_batchnorm': 'before', 'dense_batchnorm': 'after', 'conv_dropout': 0.19179270727092163, 'dense_dropout': 0.25768620210213067}. Best is trial 0 with value: 160.976891040802.

Epoch:	10	Training Loss:	4.172	Validation Loss:	4.024
--------	----	----------------	-------	------------------	-------

Files already downloaded and verified

Epoch:	1	Training Loss:	4.667	Validation Loss:	4.611
Epoch:	2	Training Loss:	4.628	Validation Loss:	4.510
Epoch:	3	Training Loss:	4.547	Validation Loss:	4.428
Epoch:	4	Training Loss:	4.485	Validation Loss:	4.387
Epoch:	5	Training Loss:	4.445	Validation Loss:	4.362
Epoch:	6	Training Loss:	4.424	Validation Loss:	4.356
Epoch:	7	Training Loss:	4.407	Validation Loss:	4.330
Epoch:	8	Training Loss:	4.395	Validation Loss:	4.330
Epoch:	9	Training Loss:	4.377	Validation Loss:	4.362

[I 2022-01-24 03:55:07,186] Trial 1 finished with value: 171.43050575256348 and parameters: {'optimizer': 'adam', 'num_conv_blocks': 4, 'num_dense_layers': 2, 'conv_kernel_size': 2, 'pool_kernel_size': 5, 'conv_activations': 'sigmoid', 'dense_activations': 'leakyrelu', 'conv_batchnorm': 'before', 'dense_batchnorm': 'before', 'conv_dropout': 0.4005373293909916, 'dense_dropout': 0.12951364524530612}. Best is trial 0 with value: 160.976891040802.

Epoch:	10	Training Loss:	4.350	Validation Loss:	4.286
--------	----	----------------	-------	------------------	-------

Files already downloaded and verified

Epoch:	1	Training Loss:	4.657	Validation Loss:	4.359
Epoch:	2	Training Loss:	4.520	Validation Loss:	4.283
Epoch:	3	Training Loss:	4.457	Validation Loss:	4.221
Epoch:	4	Training Loss:	4.417	Validation Loss:	4.192
Epoch:	5	Training Loss:	4.373	Validation Loss:	4.179

Epoch: 6	Training Loss: 4.348	Validation Loss: 4.124
Epoch: 7	Training Loss: 4.326	Validation Loss: 4.100
Epoch: 8	Training Loss: 4.298	Validation Loss: 4.081
Epoch: 9	Training Loss: 4.287	Validation Loss: 4.062

[I 2022-01-24 04:10:04,541] Trial 2 finished with value: 162.40307450294495 and parameters: {'optimizer': 'sgd', 'num_conv_blocks': 2, 'num_dense_layers': 2, 'conv_kernel_size': 5, 'pool_kernel_size': 2, 'conv_activations': 'tanh', 'dense_activations': 'tanh', 'conv_batchnorm': 'after', 'dense_batchnorm': 'after', 'conv_dropout': 0.4945721582411171, 'dense_dropout': 0.3600449294674264}. Best is trial 0 with value: 160.976891040802.

Epoch: 10	Training Loss: 4.265	Validation Loss: 4.060
Files already downloaded and verified		
Epoch: 1	Training Loss: 4.657	Validation Loss: 4.467
Epoch: 2	Training Loss: 4.554	Validation Loss: 4.378
Epoch: 3	Training Loss: 4.483	Validation Loss: 4.323
Epoch: 4	Training Loss: 4.425	Validation Loss: 4.269
Epoch: 5	Training Loss: 4.385	Validation Loss: 4.225
Epoch: 6	Training Loss: 4.348	Validation Loss: 4.194
Epoch: 7	Training Loss: 4.313	Validation Loss: 4.181
Epoch: 8	Training Loss: 4.289	Validation Loss: 4.124
Epoch: 9	Training Loss: 4.263	Validation Loss: 4.107

[I 2022-01-24 04:29:34,390] Trial 3 finished with value: 163.596444606781 and parameters: {'optimizer': 'sgd', 'num_conv_blocks': 4, 'num_dense_layers': 3, 'conv_kernel_size': 8, 'pool_kernel_size': 3, 'conv_activations': 'tanh', 'dense_activations': 'leakyrelu', 'conv_batchnorm': 'before', 'dense_batchnorm': 'before', 'conv_dropout': 0.12940670912374097, 'dense_dropout': 0.44156351204304856}. Best is trial 0 with value: 160.976891040802.

Epoch: 10	Training Loss: 4.241	Validation Loss: 4.090
Files already downloaded and verified		
Epoch: 1	Training Loss: 4.611	Validation Loss: 4.347
Epoch: 2	Training Loss: 4.466	Validation Loss: 4.303
Epoch: 3	Training Loss: 4.411	Validation Loss: 4.263
Epoch: 4	Training Loss: 4.372	Validation Loss: 4.198
Epoch: 5	Training Loss: 4.334	Validation Loss: 4.168
Epoch: 6	Training Loss: 4.317	Validation Loss: 4.127
Epoch: 7	Training Loss: 4.296	Validation Loss: 4.145
Epoch: 8	Training Loss: 4.286	Validation Loss: 4.138
Epoch: 9	Training Loss: 4.272	Validation Loss: 4.077

[I 2022-01-24 04:45:15,816] Trial 4 finished with value: 163.43482208251953 and parameters: {'optimizer': 'adam', 'num_conv_blocks': 3, 'num_dense_layers': 4, 'conv_kernel_size': 2, 'pool_kernel_size': 4, 'conv_activations': 'relu', 'dense_activations': 'tanh', 'conv_batchnorm': 'before', 'dense_batchnorm': 'after', 'conv_dropout': 0.3430825134068658,

'dense_dropout': 0.21658675460165494}. Best is trial 0 with value: 160.976891040802.

Epoch: 10	Training Loss: 4.262	Validation Loss: 4.086
-----------	----------------------	------------------------

Files already downloaded and verified

Epoch: 1	Training Loss: 4.370	Validation Loss: 4.137
Epoch: 2	Training Loss: 4.197	Validation Loss: 4.026
Epoch: 3	Training Loss: 4.119	Validation Loss: 3.950
Epoch: 4	Training Loss: 4.071	Validation Loss: 3.899
Epoch: 5	Training Loss: 4.032	Validation Loss: 3.854
Epoch: 6	Training Loss: 3.994	Validation Loss: 3.795
Epoch: 7	Training Loss: 3.966	Validation Loss: 3.783
Epoch: 8	Training Loss: 3.938	Validation Loss: 3.760
Epoch: 9	Training Loss: 3.911	Validation Loss: 3.724

[I 2022-01-24 05:00:24,257] Trial 5 finished with value: 148.07879781723022 and parameters: {'optimizer': 'adam', 'num_conv_blocks': 2, 'num_dense_layers': 1, 'conv_kernel_size': 9, 'pool_kernel_size': 2, 'conv_activations': 'relu', 'dense_activations': 'leakyrelu', 'conv_batchnorm': 'after', 'dense_batchnorm': 'before', 'conv_dropout': 0.41786473811059455, 'dense_dropout': 0.16982388420062394}. Best is trial 5 with value: 148.07879781723022.

Epoch: 10	Training Loss: 3.894	Validation Loss: 3.702
-----------	----------------------	------------------------

Files already downloaded and verified

Epoch: 1	Training Loss: 4.674	Validation Loss: 4.626
Epoch: 2	Training Loss: 4.651	Validation Loss: 4.614
Epoch: 3	Training Loss: 4.644	Validation Loss: 4.606
Epoch: 4	Training Loss: 4.637	Validation Loss: 4.602
Epoch: 5	Training Loss: 4.637	Validation Loss: 4.599
Epoch: 6	Training Loss: 4.631	Validation Loss: 4.597
Epoch: 7	Training Loss: 4.631	Validation Loss: 4.594
Epoch: 8	Training Loss: 4.628	Validation Loss: 4.590
Epoch: 9	Training Loss: 4.629	Validation Loss: 4.587

[I 2022-01-24 05:15:32,249] Trial 6 finished with value: 183.25664234161377 and parameters: {'optimizer': 'sgd', 'num_conv_blocks': 2, 'num_dense_layers': 2, 'conv_kernel_size': 8, 'pool_kernel_size': 3, 'conv_activations': 'sigmoid', 'dense_activations': 'sigmoid', 'conv_batchnorm': 'before', 'dense_batchnorm': 'before', 'conv_dropout': 0.4962010530402665, 'dense_dropout': 0.3713930295322787}. Best is trial 5 with value: 148.07879781723022.

Epoch: 10	Training Loss: 4.623	Validation Loss: 4.581
-----------	----------------------	------------------------

Files already downloaded and verified

Epoch: 1	Training Loss: 4.581	Validation Loss: 4.283
Epoch: 2	Training Loss: 4.339	Validation Loss: 4.136
Epoch: 3	Training Loss: 4.235	Validation Loss: 4.068
Epoch: 4	Training Loss: 4.180	Validation Loss: 3.992
Epoch: 5	Training Loss: 4.137	Validation Loss: 3.953

Epoch: 6	Training Loss: 4.099	Validation Loss: 3.913
Epoch: 7	Training Loss: 4.071	Validation Loss: 3.878
Epoch: 8	Training Loss: 4.047	Validation Loss: 3.841
Epoch: 9	Training Loss: 4.031	Validation Loss: 3.812

[I 2022-01-24 05:31:27,581] Trial 7 finished with value: 152.01260638237 and parameters: {'optimizer': 'adam', 'num_conv_blocks': 4, 'num_dense_layers': 1, 'conv_kernel_size': 3, 'pool_kernel_size': 5, 'conv_activations': 'relu', 'dense_activations': 'leakyrelu', 'conv_batchnorm': 'after', 'dense_batchnorm': 'after', 'conv_dropout': 0.4260640048050487, 'dense_dropout': 0.3556335650272967}. Best is trial 5 with value: 148.07879781723022.

Epoch: 10	Training Loss: 4.008	Validation Loss: 3.800
Files already downloaded and verified		
Epoch: 1	Training Loss: 4.494	Validation Loss: 4.186
Epoch: 2	Training Loss: 4.301	Validation Loss: 4.124
Epoch: 3	Training Loss: 4.243	Validation Loss: 4.094
Epoch: 4	Training Loss: 4.201	Validation Loss: 4.067
Epoch: 5	Training Loss: 4.179	Validation Loss: 4.045
Epoch: 6	Training Loss: 4.157	Validation Loss: 4.019
Epoch: 7	Training Loss: 4.144	Validation Loss: 4.027
Epoch: 8	Training Loss: 4.145	Validation Loss: 4.001
Epoch: 9	Training Loss: 4.122	Validation Loss: 3.987

[I 2022-01-24 05:47:01,481] Trial 8 finished with value: 160.26873660087585 and parameters: {'optimizer': 'adam', 'num_conv_blocks': 3, 'num_dense_layers': 3, 'conv_kernel_size': 6, 'pool_kernel_size': 2, 'conv_activations': 'leakyrelu', 'dense_activations': 'sigmoid', 'conv_batchnorm': 'before', 'dense_batchnorm': 'after', 'conv_dropout': 0.13013382314348784, 'dense_dropout': 0.40888582990951783}. Best is trial 5 with value: 148.07879781723022.

Epoch: 10	Training Loss: 4.109	Validation Loss: 4.007
Files already downloaded and verified		
Epoch: 1	Training Loss: 4.619	Validation Loss: 4.452
Epoch: 2	Training Loss: 4.509	Validation Loss: 4.364
Epoch: 3	Training Loss: 4.445	Validation Loss: 4.305
Epoch: 4	Training Loss: 4.387	Validation Loss: 4.252
Epoch: 5	Training Loss: 4.347	Validation Loss: 4.212
Epoch: 6	Training Loss: 4.311	Validation Loss: 4.189
Epoch: 7	Training Loss: 4.285	Validation Loss: 4.156
Epoch: 8	Training Loss: 4.258	Validation Loss: 4.118
Epoch: 9	Training Loss: 4.229	Validation Loss: 4.091

[I 2022-01-24 06:02:53,443] Trial 9 finished with value: 163.08367204666138 and parameters: {'optimizer': 'sgd', 'num_conv_blocks': 3, 'num_dense_layers': 4, 'conv_kernel_size': 4, 'pool_kernel_size': 2, 'conv_activations': 'tanh', 'dense_activations': 'leakyrelu', 'conv_batchnorm': 'after', 'dense_batchnorm': 'before', 'conv_dropout': 0.3574359087360288,

'dense_dropout': 0.22527667617198507}. Best is trial 5 with value: 148.07879781723022.

Epoch: 10 |Training Loss: 4.210 |Validation Loss: 4.077

HPO process took 9691 seconds.

Number of completed trials: 10

Best trial:

```
FrozenTrial(number=5, values=[148.07879781723022],
datetime_start=datetime.datetime(2022, 1, 24, 4, 45, 15, 817464),
datetime_complete=datetime.datetime(2022, 1, 24, 5, 0, 24, 256433),
params={'optimizer': 'adam', 'num_conv_blocks': 2, 'num_dense_layers': 1,
'conv_kernel_size': 9, 'pool_kernel_size': 2, 'conv_activations': 'relu',
'dense_activations': 'leakyrelu', 'conv_batchnorm': 'after', 'dense_batchnorm':
'before', 'conv_dropout': 0.41786473811059455, 'dense_dropout':
0.16982388420062394}, distributions={'optimizer':
CategoricalDistribution(choices=('adam', 'sgd')), 'num_conv_blocks':
IntUniformDistribution(high=5, low=2, step=1), 'num_dense_layers':
IntUniformDistribution(high=5, low=1, step=1), 'conv_kernel_size':
IntUniformDistribution(high=10, low=2, step=1), 'pool_kernel_size':
IntUniformDistribution(high=5, low=2, step=1), 'conv_activations':
CategoricalDistribution(choices=('relu', 'leakyrelu', 'sigmoid', 'tanh')),
'dense_activations': CategoricalDistribution(choices=('relu', 'leakyrelu',
'sigmoid', 'tanh')), 'conv_batchnorm':
CategoricalDistribution(choices=('before', 'after')), 'dense_batchnorm':
CategoricalDistribution(choices=('before', 'after')), 'conv_dropout':
UniformDistribution(high=0.5, low=0.1), 'dense_dropout':
UniformDistribution(high=0.5, low=0.1)}, user_attrs={}, system_attrs={},
intermediate_values={}, trial_id=5, state=TrialState.COMPLETE, value=None)
```

Best Validation Loss: 148.07879781723022

Best Params:

```
optimizer: adam
num_conv_blocks: 2
num_dense_layers: 1
conv_kernel_size: 9
pool_kernel_size: 2
conv_activations: relu
dense_activations: leakyrelu
conv_batchnorm: after
dense_batchnorm: before
conv_dropout: 0.41786473811059455
dense_dropout: 0.16982388420062394
```

```
[ ]: # Selecting the best algorithm
torch.cuda.empty_cache()
best_hparams = trial.params
hparams.update(best_hparams)

# Training model with these parameters
```

```

hparams['epochs'] = 80
best_model = ConvNet(hparams)

inzvaNet.to(device)
inzvaNet.train()
criterion = hparams['loss_function']
if hparams['optimizer']=='adam':
    if hparams.get('optimizer_params'):
        optimizer = optim.Adam(inzvaNet.parameters(), lr=hparams['lr'],
↪**hparams['optimizer_params'])
    else:
        optimizer = optim.Adam(inzvaNet.parameters(), lr=hparams['lr'])
elif hparams['optimizer']=='sgd':
    if hparams.get('optimizer_params'):
        optimizer = optim.SGD(inzvaNet.parameters(), lr=hparams['lr'],
↪**hparams['optimizer_params'])
    else:
        optimizer = optim.SGD(inzvaNet.parameters(), lr=hparams['lr'])
num_training_batches = len(trainloader)
num_validation_batches = len(validloader)

print('training starts!')
num_training_batches = len(trainloader)
num_validation_batches = len(validloader)
num_testing_batches = len(testloader)

tStart = timer()
for epoch in range(hparams["epochs"]):

    epoch_loss_training = 0.0
    epoch_loss_validation = 0.0

    inzvaNet.train()
    for i, data in enumerate(trainloader):
        images, labels = data[0].to(device), data[1].to(device)
        optimizer.zero_grad()
        predictions = inzvaNet(images)
        loss = criterion(predictions, labels)
        loss.backward()
        optimizer.step()
        epoch_loss_training += loss.item()

    inzvaNet.eval()
    with torch.no_grad():
        for i, data in enumerate(validloader):
            images, labels = data[0].to(device), data[1].to(device)
            predictions = inzvaNet(images)

```

```

        loss = criterion(predictions, labels)
        epoch_loss_validation += loss.item()

    print("Epoch: %2d\t|Training Loss: %.3f\t|Validation Loss: %.3f" %
    ↪(epoch + 1,
        epoch_loss_training / num_training_batches,
        epoch_loss_validation / num_validation_batches))

tFinish = timer()
print('Finished Training.')
print("Training process took %.2f seconds"%(tFinish-tStart))
print("Saving model...")
try:
    torch.save(inzvaNet, "ConvNet_best.pth")
except Exception as e:
    print(e)
    print("Failed to save the model.")
print("Done saving.")

# Evaluation of the model on testset
inzvaNet.eval()
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = inzvaNet(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print("correct predictions: ", correct)
print("total number of data: ", total)
print('Accuracy of the network on all of the test images: %d %%' % (100 *
    ↪correct / total))

# Class-wise accuracy
class_correct = list(0. for i in range(100))
class_total = list(0. for i in range(100))
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = inzvaNet(images)
        _, predicted = torch.max(outputs, 1)
        c = (predicted == labels).squeeze()

```

```

    for i in range(len(labels)):
        label = labels[i]
        class_correct[label] += c[i].item()
        class_total[label] += 1

print("Accuracy of classes:")
for i in range(100):
    print('%3d- %20s : %5.3f %%' % (
        i+1, classes[i], 100 * class_correct[i] / class_total[i]))

```

training starts!

Epoch: 1	Training Loss:	3.808	Validation Loss:	3.634
Epoch: 2	Training Loss:	3.790	Validation Loss:	3.617
Epoch: 3	Training Loss:	3.776	Validation Loss:	3.596
Epoch: 4	Training Loss:	3.752	Validation Loss:	3.581
Epoch: 5	Training Loss:	3.733	Validation Loss:	3.582
Epoch: 6	Training Loss:	3.720	Validation Loss:	3.579
Epoch: 7	Training Loss:	3.709	Validation Loss:	3.541
Epoch: 8	Training Loss:	3.705	Validation Loss:	3.539
Epoch: 9	Training Loss:	3.690	Validation Loss:	3.539
Epoch: 10	Training Loss:	3.684	Validation Loss:	3.532
Epoch: 11	Training Loss:	3.665	Validation Loss:	3.492
Epoch: 12	Training Loss:	3.667	Validation Loss:	3.477
Epoch: 13	Training Loss:	3.651	Validation Loss:	3.494
Epoch: 14	Training Loss:	3.648	Validation Loss:	3.491
Epoch: 15	Training Loss:	3.638	Validation Loss:	3.463
Epoch: 16	Training Loss:	3.640	Validation Loss:	3.478
Epoch: 17	Training Loss:	3.626	Validation Loss:	3.437
Epoch: 18	Training Loss:	3.619	Validation Loss:	3.435
Epoch: 19	Training Loss:	3.603	Validation Loss:	3.407
Epoch: 20	Training Loss:	3.613	Validation Loss:	3.434
Epoch: 21	Training Loss:	3.605	Validation Loss:	3.439
Epoch: 22	Training Loss:	3.596	Validation Loss:	3.405
Epoch: 23	Training Loss:	3.594	Validation Loss:	3.408
Epoch: 24	Training Loss:	3.585	Validation Loss:	3.438
Epoch: 25	Training Loss:	3.573	Validation Loss:	3.402
Epoch: 26	Training Loss:	3.576	Validation Loss:	3.383
Epoch: 27	Training Loss:	3.575	Validation Loss:	3.399
Epoch: 28	Training Loss:	3.569	Validation Loss:	3.372
Epoch: 29	Training Loss:	3.561	Validation Loss:	3.379
Epoch: 30	Training Loss:	3.561	Validation Loss:	3.369
Epoch: 31	Training Loss:	3.557	Validation Loss:	3.344
Epoch: 32	Training Loss:	3.549	Validation Loss:	3.370
Epoch: 33	Training Loss:	3.550	Validation Loss:	3.359
Epoch: 34	Training Loss:	3.537	Validation Loss:	3.332
Epoch: 35	Training Loss:	3.540	Validation Loss:	3.324
Epoch: 36	Training Loss:	3.535	Validation Loss:	3.320
Epoch: 37	Training Loss:	3.520	Validation Loss:	3.321

Epoch: 38	Training Loss:	3.526	Validation Loss:	3.360
Epoch: 39	Training Loss:	3.524	Validation Loss:	3.319
Epoch: 40	Training Loss:	3.520	Validation Loss:	3.308
Epoch: 41	Training Loss:	3.513	Validation Loss:	3.330
Epoch: 42	Training Loss:	3.513	Validation Loss:	3.331
Epoch: 43	Training Loss:	3.509	Validation Loss:	3.326
Epoch: 44	Training Loss:	3.509	Validation Loss:	3.311
Epoch: 45	Training Loss:	3.505	Validation Loss:	3.321
Epoch: 46	Training Loss:	3.502	Validation Loss:	3.280
Epoch: 47	Training Loss:	3.500	Validation Loss:	3.314
Epoch: 48	Training Loss:	3.500	Validation Loss:	3.302
Epoch: 49	Training Loss:	3.500	Validation Loss:	3.291
Epoch: 50	Training Loss:	3.488	Validation Loss:	3.289
Epoch: 51	Training Loss:	3.486	Validation Loss:	3.271
Epoch: 52	Training Loss:	3.480	Validation Loss:	3.311
Epoch: 53	Training Loss:	3.480	Validation Loss:	3.271
Epoch: 54	Training Loss:	3.476	Validation Loss:	3.285
Epoch: 55	Training Loss:	3.476	Validation Loss:	3.287
Epoch: 56	Training Loss:	3.470	Validation Loss:	3.270
Epoch: 57	Training Loss:	3.481	Validation Loss:	3.256
Epoch: 58	Training Loss:	3.466	Validation Loss:	3.250
Epoch: 59	Training Loss:	3.473	Validation Loss:	3.284
Epoch: 60	Training Loss:	3.477	Validation Loss:	3.249
Epoch: 61	Training Loss:	3.474	Validation Loss:	3.262
Epoch: 62	Training Loss:	3.462	Validation Loss:	3.260
Epoch: 63	Training Loss:	3.457	Validation Loss:	3.261
Epoch: 64	Training Loss:	3.460	Validation Loss:	3.246
Epoch: 65	Training Loss:	3.455	Validation Loss:	3.266
Epoch: 66	Training Loss:	3.455	Validation Loss:	3.245
Epoch: 67	Training Loss:	3.451	Validation Loss:	3.242
Epoch: 68	Training Loss:	3.446	Validation Loss:	3.237
Epoch: 69	Training Loss:	3.448	Validation Loss:	3.235
Epoch: 70	Training Loss:	3.444	Validation Loss:	3.236
Epoch: 71	Training Loss:	3.448	Validation Loss:	3.236
Epoch: 72	Training Loss:	3.438	Validation Loss:	3.243
Epoch: 73	Training Loss:	3.436	Validation Loss:	3.223
Epoch: 74	Training Loss:	3.440	Validation Loss:	3.198
Epoch: 75	Training Loss:	3.436	Validation Loss:	3.228
Epoch: 76	Training Loss:	3.440	Validation Loss:	3.212
Epoch: 77	Training Loss:	3.442	Validation Loss:	3.224
Epoch: 78	Training Loss:	3.429	Validation Loss:	3.195
Epoch: 79	Training Loss:	3.436	Validation Loss:	3.239
Epoch: 80	Training Loss:	3.436	Validation Loss:	3.207

Finished Training.

Training process took 7008.78 seconds.

Saving model...

Done saving.

correct predictions: 2151

total number of data: 10000

Accuracy of the network on all of the test images: 21 %

Accuracy of classes:

1-	apples	: 35.000 %
2-	aquarium fish	: 18.000 %
3-	baby	: 15.000 %
4-	bear	: 10.000 %
5-	beaver	: 10.000 %
6-	bed	: 20.000 %
7-	bee	: 20.000 %
8-	beetle	: 31.000 %
9-	bicycle	: 22.000 %
10-	bottles	: 8.000 %
11-	bowls	: 0.000 %
12-	boy	: 1.000 %
13-	bridge	: 24.000 %
14-	bus	: 17.000 %
15-	butterfly	: 5.000 %
16-	camel	: 18.000 %
17-	cans	: 16.000 %
18-	castle	: 37.000 %
19-	caterpillar	: 25.000 %
20-	cattle	: 5.000 %
21-	chair	: 37.000 %
22-	chimpanzee	: 47.000 %
23-	clock	: 12.000 %
24-	cloud	: 64.000 %
25-	cockroach	: 60.000 %
26-	computer keyboard	: 4.000 %
27-	couch	: 4.000 %
28-	crab	: 9.000 %
29-	crocodile	: 3.000 %
30-	cups	: 7.000 %
31-	dinosaur	: 34.000 %
32-	dolphin	: 23.000 %
33-	elephant	: 2.000 %
34-	flatfish	: 25.000 %
35-	forest	: 12.000 %
36-	fox	: 14.000 %
37-	girl	: 26.000 %
38-	hamster	: 9.000 %
39-	house	: 14.000 %
40-	kangaroo	: 21.000 %
41-	lamp	: 1.000 %
42-	lawn-mower	: 43.000 %
43-	leopard	: 24.000 %
44-	lion	: 53.000 %
45-	lizard	: 12.000 %

46-	lobster	: 20.000 %
47-	man	: 8.000 %
48-	maple	: 19.000 %
49-	motorcycle	: 52.000 %
50-	mountain	: 14.000 %
51-	mouse	: 0.000 %
52-	mushrooms	: 9.000 %
53-	oak	: 69.000 %
54-	oranges	: 49.000 %
55-	orchids	: 49.000 %
56-	otter	: 1.000 %
57-	palm	: 36.000 %
58-	pears	: 12.000 %
59-	pickup truck	: 13.000 %
60-	pine	: 9.000 %
61-	plain	: 54.000 %
62-	plates	: 23.000 %
63-	poppies	: 39.000 %
64-	porcupine	: 27.000 %
65-	possum	: 6.000 %
66-	rabbit	: 1.000 %
67-	raccoon	: 22.000 %
68-	ray	: 20.000 %
69-	road	: 59.000 %
70-	rocket	: 28.000 %
71-	roses	: 24.000 %
72-	sea	: 33.000 %
73-	seal	: 4.000 %
74-	shark	: 33.000 %
75-	shrew	: 9.000 %
76-	skunk	: 45.000 %
77-	skyscraper	: 37.000 %
78-	snail	: 10.000 %
79-	snake	: 6.000 %
80-	spider	: 24.000 %
81-	squirrel	: 6.000 %
82-	streetcar	: 7.000 %
83-	sunflowers	: 69.000 %
84-	sweet peppers	: 22.000 %
85-	table	: 5.000 %
86-	tank	: 19.000 %
87-	telephone	: 29.000 %
88-	television	: 10.000 %
89-	tiger	: 46.000 %
90-	tractor	: 41.000 %
91-	train	: 18.000 %
92-	trout	: 13.000 %
93-	tulips	: 8.000 %

94-	turtle	:	7.000	%
95-	wardrobe	:	45.000	%
96-	whale	:	39.000	%
97-	willow	:	29.000	%
98-	wolf	:	18.000	%
99-	woman	:	7.000	%
100-	worm	:	7.000	%